Table 1: OGBench (proprioceptive input), success rate (%). For FB, HILP and RLDP the "paper" column is the published number from Table 1 of TD-JEPA and the "ours" column is our reimplementation evaluated by the same reward-inference route. Superscripts mark the Density-FB evaluation route: $^{\mathrm{G}}$ goal-conditioned, $^{\mathrm{R}}$ reward-inference.

	FB		HILP		RLDP		Density-FB (ours)	
	paper	ours	paper	ours	paper	ours	goal$^{\mathrm{G}}$	z^{R}
avg	39.0 ±0.7	—	38.0 ±1.1	—	27.1 ±0.8	—	—	—
antmaze-mn	73.0 ±2.7	—	83.6 ±2.6	—	74.6 ±4.2	—	66.9 ±24.0	—
antmaze-ln	36.8 ±4.3	—	52.6 ±3.9	—	36.4 ±4.7	—	—	—
antmaze-ms	70.4 ±4.0	—	50.6 ±2.5	—	58.4 ±3.3	—	—	—
antmaze-ls	49.8 ±5.6	—	12.2 ±1.8	—	19.6 ±2.7	—	—	—
antmaze-me	51.6 ±2.6	—	2.0 ±0.8	—	4.8 ±2.4	—	—	—
cube-single	49.6 ±3.8	39.0	74.2 ±3.5	83.0	19.8 ±2.4	25.0	91.5 ±9.8	92.0
cube-double	2.6 ±0.4	—	20.0 ±2.7	8.0	3.8 ±0.8	—	68.0 ±7.0	20.0
scene	12.8 ±1.6	—	43.8 ±1.9	—	11.6 ±1.6	—	64.1 ±10.9	—
puzzle-3x3	4.8 ±0.7	—	2.8 ±0.7	—	14.6 ±0.9	—	—	—

Paper columns: published numbers, 1M steps, 10 seeds, best cell of a 4-cell hyperparameter grid, reward-inference eval (z inferred from 10k relabelled samples); mean ± s.e. Ours columns: our reimplementations, state input, flow actor, same reward-inference eval, 1M steps, 3 seeds (vs. the paper's 10 seeds with best-of-4 grid), at the TD-JEPA OGBench sweep cell FB ortho = 1000 / $\mathrm{lr}_B = 10^{-4}$, HILP ϕ-lr = 10^{-4} / $p_{\text{rand-goal}} = 0.375$, RLDP ortho = 0.01 / $\mathrm{lr}_B = 10^{-4}$; the FB cube-single cell mixes $\mathrm{lr}_B = 10^{-4}$ and 10^{-5} runs (re-run pending). Density-FB$^{\mathrm{G}}$: goal-conditioned eval (z from the goal observation), state input, 500k steps, 3 seeds; mean ± std over seeds. Density-FB$^{\mathrm{R}}$: reward-inference eval (z from 10k relabelled samples), state input, 1M steps; cube-single 3 seeds, cube-double 1 seed. "—": not run.