Table 1: Does our reimplementation reproduce the published baselines? OGBench manipulation (proprioceptive input), success rate (%), 1M steps, reward-inference evaluation. "published" is TD-JEPA Table 1; "ours" is our code at the same hyperparameter cell, under the protocol every number in this paper uses. RLDP is also shown at its own default orthonormality coefficient 1.0, which TD-JEPA's grid did not contain.

	FB		HILP		RLDP		
	published	ours	published	ours	published	ours (ortho 0.01)	ours (ortho 1)
avg	17.5	15.2	35.2	30.7	12.5	8.6	12.7
cube-single	49.6 ±3.8	45.6 ±5.9	74.2 ±3.5	71.1 ±5.7	19.8 ±2.4	25.0	27.6 ±3.7
cube-double	2.6 ±0.4	2.2 ±2.7	20.0 ±2.7	11.1 ±3.3	3.8 ±0.8	2.2 ±1.1	2.7 ±1.3
scene	12.8 ±1.6	10.2 ±4.1	43.8 ±1.9	39.3 ±5.7	11.6 ±1.6	4.7 ±1.1	10.2 ±3.2
puzzle-3x3	4.8 ±0.7	2.7 ±0.9	2.8 ±0.7	1.3 ±0.5	14.6 ±0.9	2.4 ±1.1	10.4 ±1.7

Published columns: 1M steps, 10 seeds, best cell of a 4-cell hyperparameter grid, reward-inference eval (z inferred from 10^4 relabelled samples); mean ± s.e. The five antmaze rows of that table are dropped: this paper is about manipulation, and the object variable that motivates the method does not exist in navigation. Ours columns: state input, flow actor, same reward-inference eval, last-3-checkpoint estimator, 1M steps, 3 seeds, at the TD-JEPA OGBench sweep cell — FB ortho = 1000 / $\mathrm{lr}_B = 10^{-4}$, HILP ϕ-lr = 10^{-4} / $p_{\text{rand-goal}} = 0.375$, RLDP ortho = 0.01 / $\mathrm{lr}_B = 10^{-4}$; mean ± std. FB and HILP cube cells come from the disjoint 3-site campaign (each cell 3 independent seeds, no shared save-dir). RLDP cube-single at ortho 0.01 is an eval-only rescore with a single value on file. RLDP at ortho 1 (exp/ortho1) is the configuration used everywhere else in this paper: it is RLDP's own published default, and it is worth scene 4.7 → 10.2 and puzzle 2.4 → 10.4 over TD-JEPA's grid value. Higher coefficients than 1 are worse on every task. Reading: FB and HILP reproduce to within a few points on every task (FB 15.2 vs 17.5 on average, HILP 30.7 vs 35.2); RLDP reproduces only once its own orthonormality coefficient is restored (12.7 vs 12.5), and puzzle-3x3 is where our RLDP stays furthest below the published number (10.4 vs 14.6). **avg** is the mean of the four rows, recomputed per column; no spread is carried through the recomputation. On cube-double, reward inference is capped near half of goal-conditioned success for every method: OGBench re-permutes the cube-to-target assignment at every reset, while the reward samples used to infer z encode one fixed assignment. The density objective inside each of these backbones is Table **??**; every retrofit variant is Table **??**.