Table 1: OGBench (proprioceptive input), success rate (%). Baseline columns are the published numbers from Table 1 of TD-JEPA; the two right-hand columns are Density-FB. Superscripts mark the evaluation route: $^{\mathrm{G}}$ goal-conditioned, $^{\mathrm{R}}$ reward-inference.

	published			Density-FB (ours)	
	FB	HILP	RLDP	goal$^{\mathrm{G}}$	z^{R}
avg	39.0 ±0.7	38.0 ±1.1	27.1 ±0.8	—	—
antmaze-mn	73.0 ±2.7	83.6 ±2.6	74.6 ±4.2	66.9 ±24.0	—
antmaze-ln	36.8 ±4.3	52.6 ±3.9	36.4 ±4.7	—	—
antmaze-ms	70.4 ±4.0	50.6 ±2.5	58.4 ±3.3	—	—
antmaze-ls	49.8 ±5.6	12.2 ±1.8	19.6 ±2.7	—	—
antmaze-me	51.6 ±2.6	2.0 ±0.8	4.8 ±2.4	—	—
cube-single	49.6 ±3.8	74.2 ±3.5	19.8 ±2.4	91.5 ±9.8	92.0
cube-double	2.6 ±0.4	20.0 ±2.7	3.8 ±0.8	68.0 ±7.0	20.0
scene	12.8 ±1.6	43.8 ±1.9	11.6 ±1.6	64.1 ±10.9	—
puzzle-3x3	4.8 ±0.7	2.8 ±0.7	14.6 ±0.9	—	—

Baselines (FB, HILP, RLDP): published numbers, 1M steps, 10 seeds, best cell of the hyperparameter grid, reward-inference eval (z inferred from 10k relabelled samples); mean ± s.e. Density-FB$^{\mathrm{G}}$: goal-conditioned eval (z from the goal observation), state input, 500k steps, 3 seeds; mean ± std over seeds. Density-FB$^{\mathrm{R}}$: reward-inference eval (z from 10k relabelled samples), state input, 1M steps; cube-single 3 seeds, cube-double 1 seed. "—": not run.