Table 1: The density objective inside each backbone without privileged input, and the object-coordinate control. OGBench `cube-single`, success rate (%), 1M steps, reward-inference evaluation (goal conditioning in the second column). "B input": what the backward / feature network receives. "tasks": feature = task vector from a random dataset state's features (the host method's rule); regression = hindsight goal from the same trajectory, task vector by the reward regression of Appendix **??**. RLDP rows use its default orthonormality coefficient 1.0. Mean ± std over 3 seeds.

Critic	B input	tasks	CE trains B	reward inference	goal conditioning
Full state (no privileged input)					
FB	state	feature	–	45.6 ±5.9	—
Density-FB	state	regression	yes	90.9	84.3 ±5.1[f]
RLDP	state	feature	–	27.7 ±3.5	15.3 ±2.1
Density-RLDP	state	feature	no	36.7 ±2.9	13.3 ±3.8
Density-RLDP	state	feature	yes	32.3 ±5.5	17.0 ±3.0
Density-RLDP	state	regression	no	53.3 ±3.2	50.7 ±4.0
HILP	state	feature	–	71.1 ±5.7	—
Density-HILP	state	feature	no	71.1 ±4.8	—
Density-HILP	state	regression	no	77.0 ±3.5	70.3 ±2.3
Object-coordinate B input (privileged)					
Density-FB (Table **??**)	object coords	regression	yes	81.3 ±3.8	74.9 ±5.8
RLDP	object coords	feature	–	87.7 ±4.0	81.0 ±2.0
Density-RLDP	object coords	feature	no	87.7 ±4.0	89.3 ±2.9
Density-RLDP	object coords	regression	yes	79.7 ±9.0	73.7 ±17.4
HILP	object coords	feature	–	running	
Density-HILP	object coords	regression	no	queued	

All rows: state observations to the forward network and actor, flow actor, reward-inference evaluation with 10^4 relabelled samples, last-3-checkpoint estimator, 1M steps, 3 seeds. Sources: FB / HILP / Density-HILP(feature) from the campaign (hub); Density-FB state-only = density_fb_z_state (hub, vista); RLDP and Density-RLDP(feature) at ortho 1 = exp/ortho1; Density-RLDP levers = exp/levers; Density-HILP regression, RLDP object coords, HILP object coords = exp/hlevers. Density-FB state: reward inference is the campaign last-3 aggregate (no per-seed spread on file), goal route [f] = final-step mean of 80/90/83. Per-seed reward-inference values: RLDP 24/31/27; Density-RLDP 35/35/40; CE-trains-B 26/36/35; regression 52/51/57; Density-HILP regression 75/75/81; RLDP object coords 84/92/87; Density-RLDP object coords 83/90/90; all three 71/79/89. Reading: without privileged input the objective improves every backbone on this task (FB 46 → 91, RLDP 28 → 53 with the regression-task lever, HILP 71 → 77). With object-coordinate input, plain RLDP already reaches 88, level with Density-RLDP, so the object-coordinate rows measure the input, not the objective.