Table 1: **Privileged object coordinates, every estimator family.** The backward tower (B for FB/RLDP, ϕ for HILP) consumes the ϵ-scaled object-coordinate block — 3 numbers (cube xyz) on `cube-single`, 6 on `cube-double` — while the forward map, the left encoder and the policy keep the whole observation. Each family is a matched pair: the vanilla estimator and the same estimator trained with the normalized categorical (density) objective. Superscripts give the number of seeds; mean ± sd, and n=1 cells should not be quoted. *Protocol:* 500k gradient steps, flow (FQL+BC) actor, 20 eval episodes per task, mean over the environment's 5 evaluation tasks — NOT the 50-episode protocol of Table **??**.

	FB		RLDP		HILP	
Dataset	vanilla	+ density	vanilla	+ density	vanilla	+ density
OGBench, object-coordinate backward map						
`cube-single`	0.340^{1}	$\mathbf{0.943 \pm 0.006}^{\mathbf{3}}$	0.530 ± 0.354^{2}	0.820 ± 0.057^{2}	0.090^{1}	0.145 ± 0.007^{2}
`cube-double`	—	$\mathbf{0.640 \pm 0.040}^{\mathbf{3}}$	0.160^{1}	0.170^{1}	0.060^{1}	0.040^{1}